

CS217 - Data Structures & Algorithm Analysis (DSAA)

Lecture #15

▶ **Minimum spanning trees and shortest paths**

Prof. Pietro S. Oliveto

Department of Computer Science and Engineering

Southern University of Science and Technology (SUSTech)

`olivetop@sustech.edu.cn`

<https://faculty.sustech.edu.cn/olivetop>

Reading: Chapter 21 and Section 22.3

► Aims for this lecture

- To introduce the minimum spanning tree problem.
- To see two different greedy approaches for solving it: Kruskal's algorithm and Prim's algorithm.
- To briefly review variants of shortest path problems.
- To cover Dijkstra's algorithm for solving the single-source shortest path problem.
 - Another example for greediness and dynamic programming
- To show how efficient data structures can be used to guarantee efficient runtimes.

► Minimum spanning trees

- Suppose we want to supply n newly built houses with electricity, using the minimum length of wire.
- Given a connected undirected graph $G = (V, E)$ where vertices represent houses (imagine one being on the grid) and edges $(u, v) \in E$ represent possible connections between houses. Each edge has a weight $w(u, v) > 0$ that gives the cost (amount of wire needed) to connect u and v .
- Looking for a subset $T \subseteq E$ of edges that connect all houses minimising the total weight $w(T) = \sum_{(u,v) \in T} w(u, v)$.
- Cycles are unhelpful, so T must be a tree!
Call it a **spanning tree** as it spans all vertices.
Looking for a **minimum(-weight) spanning tree (MST)**.

► Growing a minimum spanning tree

- Let's try to construct a minimum spanning tree iteratively by adding edges (wiring houses) to a selection $A \subseteq E$.
- This works so long as at each step the current set A is a **subset of some minimum spanning tree**.
- If we can add an edge (u, v) to A such that afterwards A is still a subset of some minimum spanning tree, the edge is called a safe edge.
 - Remember from correctness of greedy algorithms: the greedy choice is always safe.
 - We'll see how to determine which edges are safe.

► “Abstract” / Generic MST algorithm


GENERIC-MST(G, w)

```
1   $A = \emptyset$ 
2  while  $A$  does not form a spanning tree
3      find an edge  $(u, v)$  that is safe for  $A$ 
4       $A = A \cup \{(u, v)\}$ 
5  return  $A$ 
```

- Correctness of this approach by loop invariant:
 - Loop invariant: *Prior to each step, A is a subset of some MST.*
 - Initialisation: $A = \emptyset$ is a subset of some minimum spanning tree.
 - Maintenance: adding a safe edge maintains the loop invariant.
 - Termination: A is a spanning tree and a subset of an MST, so it must be an MST.
- Fair enough. But how to find a safe edge?

► Cuts

- A cut of an undirected graph $G = (V, E)$ is a partition of V in two sets $(S, V \setminus S)$.




- An edge **crosses** the cut if exactly one of its endpoints is in S .
- A cut **respects** a set A of edges if no edge in A crosses the cut.
- An edge is a **light edge** if its weight is minimal among all edges with some property, e. g. for all edges crossing the cut.

► Condition for safe edges

- **Theorem 23.1:** Let $G = (V, E)$ be a connected, undirected graph with a real-valued weight function w defined on E . Let A be a subset of E that is included in some MST for G .

If $(S, V \setminus S)$ is a cut of G that respects A , and (u, v) is a light edge crossing $(S, V \setminus S)$ then (u, v) is safe for A .

- In other words: adding a crossing edge of minimal weight to a partial MST is a **safe choice**.
- Proof is similar to the correctness of greedy algorithms, where we show that a greedy choice is safe.



► Proof of Theorem 23.1 (1)


- **Theorem 23.1:** Let $G = (V, E)$ be a connected, undirected graph with a real-valued weight function w defined on E . Let A be a subset of E that is included in some MST for G .
If $(S, V \setminus S)$ is a cut of G that respects A , and (u, v) is a light edge crossing $(S, V \setminus S)$ then (u, v) is safe for A .

Proof:

- Let T be a minimum spanning tree that includes A . A已经是最小生成树的一部分
- If T includes (u, v) , we are done. 找一个尊重A的割集，然后找个跨越割集的最短
- Now assume that T does not include (u, v) . Then we create another minimum spanning tree T' that does include (u, v) .
- We do this by cutting an edge and pasting in (u, v) .

► Proof of Theorem 23.1 (2)


- Since T is a spanning tree, the edge (u, v) forms a cycle with the simple path p from u to v in T .
- Since u and v are on different sides of the cut $(S, V - S)$, at least one edge (x, y) of p crosses the cut.
- The edge (x, y) is not in A as the cut respects A .
- Since (x, y) is on the unique simple path from u to v in T , removing (x, y) breaks T into two components.
- Adding (u, v) reconnects them to form a new spanning tree $T' = T - \{(x, y)\} \cup \{(u, v)\}$.



我替换掉的边也在同一个割集里面（并且是当前路径下唯一在割集里面的边），这个是能够安全替换的原因

► Proof of Theorem 23.1 (3)

- We show that T' is a minimum spanning tree.
- Since (u, v) is a light edge crossing $(S, V - S)$, and (x, y) also crosses the cut,
$$w(u, v) \leq w(x, y).$$
- Hence T' has weight
$$w(T') = w(T) - w(x, y) + w(u, v) \leq w(T).$$
- But T is a minimum spanning tree, hence T' must also be a minimum spanning tree.
- Why is (u, v) safe for A ? We have $A \subseteq T'$, $A \subseteq T$ (by assumption on T) and $(x, y) \notin A$, thus
$$A \cup \{(u, v)\} \subseteq T'.$$
- Adding (u, v) to A is a safe choice as we can still construct a minimum spanning tree T' .



► Edges connecting components are safe


GENERIC-MST(G, w)

```
1   $A = \emptyset$ 
2  while  $A$  does not form a spanning tree
3      find an edge  $(u, v)$  that is safe for  $A$ 
4       $A = A \cup \{(u, v)\}$ 
5  return  $A$ 
```


- The “abstract” MST algorithm (adding safe edges to A) constructs **a forest**.
- Note that initially all vertices are isolated and form their own trees.
- Theorem 23.1 implies that **for any tree T in that forest**, a light edge from T to the union of other trees is a safe edge.
 - Why? The cut $(T, V \setminus T)$ **respects the forest A** , so the theorem applies.

► Kruskal's algorithm

- Idea: **connect two trees adding an edge with minimum weight.**
- Need a way of finding out which tree a vertex belongs to.
- **Union-Find data structures** store names of sets:
 - **Find-Set(u)** returns the name of a set that element u belongs to.
 - **Union(u, v)** merges the two sets u and v belong to (if different)
 - Can be implemented efficiently with trees where the root contains the name of the set (details in Chapter 19).



(a)



(b)

► Kruskal's algorithm (2)

Ideas:

- sort all edges according to weight and process edges in this order.
- If both ends belong to different trees, add the edge and join the trees.

- **Runtime?**


- With efficient data structures for unions and finds, the runtime is dominated by the time for sorting: $O(|E|\log(|E|))$.
- Since $\log(|E|) \leq \log(|V|^2) = 2\log(|V|) = O(\log(|V|))$ we may write the runtime as $O(|E|\log(|V|))$.

KRUSKAL(G, w)

```
1:  $A = \emptyset$ 
2: for each vertex  $v \in V$  do
3:   make a set  $\{v\}$ 
4: sort the edges of  $E$  in nondecreasing order by weight  $w$ 
5: for each edge  $(u, v)$  in this order do
6:   if FIND-SET( $u$ )  $\neq$  FIND-SET( $v$ ) then
7:      $A = A \cup \{(u, v)\}$ 
8:     UNION( $u, v$ )
9: return  $A$ 
```

边的两个顶点属于不同的连通分量（不同的树），则可以安全加入生成树，并将两个连通分量合并；否则跳过这条边

► Kruskal's Algorithm: Example



► Prim's algorithm

- Alternative implementation of the “abstract” MST algorithm
- Idea: **grow a single tree** A by adding a minimum-weight edge leading away from the tree (a light edge to an isolated vertex).
- Since isolated vertices are trees, such a light edge is safe.
- How to implement Prim's algorithm efficiently?
 - Need to find a light (minimum-weight) edge to add to the tree.
 - We maintain a distance of each node to the tree (similar to BFS)
 - Initially all distances are ∞ .
 - Distances may decrease when new vertices are added to the tree.
 - Use a **Priority Queue** to keep track of the nodes with shortest distance to the current tree (light edges)

► Implementing Prim's algorithm



- Need to find a light (minimum-weight) edge to add to the tree.
- We maintain a distance “key” of each node to the tree (similar to BFS)
- Initially all distances are ∞ .
- Distances may decrease when new vertices are added to the tree.
- MST given by predecessors π (as for BFS)

PRIM(G, w, r)

```
1: for each vertex  $u \in V$  do
2:    $u.\text{key} = \infty$ 
3:    $u.\pi = \text{NIL}$ 
4:  $r.\text{key} = 0$ 
5:  $Q = V$ 
6: while  $Q \neq \emptyset$  do
7:    $u = \text{EXTRACT-MIN}(Q)$ 
8:   for each  $v \in \text{Adj}[u]$  do
9:     if  $v \in Q$  and  $w(u, v) < v.\text{key}$  then
10:       $v.\pi = u$ 
11:       $v.\text{key} = w(u, v)$ 
```

从当前生成树包含的顶点集中，找到这些顶点连接的边中权重最小的边，加入生成树

► Prim: Example



► Priority Queue based on min-heap

- A data structure for maintaining a set S of elements with an associated element called key.
- **Min-priority queue** based on min-heap defined as follows:

Operation	Time
Insert(S, x) – insert x into S	$O(\log n)$
Minimum(S) – returns smallest element in S	$O(1)$
Extract-Min(S) – removes and returns smallest element in S	$O(\log n)$
Decrease-Key(S, x, k) – decreases x 's value to smaller value k (element may float up in the heap)	$O(\log n)$

► Runtime of Prim's algorithm w/ Min-Heaps

- Runtime exclusive of red lines (as for BFS):
 $O(|V| + |E|)$
 (store a bit in each vertex to make the test $v \in Q$ run in $O(1)$ time)
- Building a Min-Heap: $O(|V|)$.
- Runtime for all calls to Extract-Min is $O(|V|\log(|V|))$. 一共 V 次, 每次 $\log V$
- Runtime for at most $|E|$ Decrease-Keys is $O(|E|\log(|V|))$.
- Total: $O(|E|\log(|V|))$** as (since G is connected) $|V| = O(|E|)$.

PRIM(G, w, r)

```

1: for each vertex  $u \in V$  do
2:    $u.\text{key} = \infty$ 
3:    $u.\pi = \text{NIL}$ 
4:  $r.\text{key} = 0$ 
5:  $Q = V$ 
6: BUILD-MIN-HEAP( $Q$ )
7: while  $Q \neq \emptyset$  do
8:    $u = \text{EXTRACT-MIN}(Q)$ 
9:   for each  $v \in \text{Adj}[u]$  do
10:    if  $v \in Q$  and  $w(u, v) < v.\text{key}$  then
11:       $v.\pi = u$ 
12:      DECREASE-KEY( $Q, v.\text{key}, w(u, v)$ )
```

每次提取堆中key最小的顶点, 更新与之相邻顶点的key

► Shortest Path Problems

- Given a **directed graph** with edge weights representing distances, what is the shortest path between two vertices?
- To find the shortest path from ShenZhen 深圳 to ShangHai 上海, exploring all paths (e.g. via Beijing 北京) is not helpful. Need a smarter approach.
- Breadth-first search finds shortest paths when all distances are 1, but can't deal with weights.
- Assume that all distances are non-negative.
- Note that shortest paths exhibit **optimal substructure**: a shortest path from s to u going through v is composed of a shortest path from s to v and a shortest path from v to u .

► Variants of Shortest Path Problems

- **Single-source shortest paths problem (SSSP)**: find shortest paths from a source vertex to all other vertices.
- **Single-destination shortest paths problem (SDSP)**: find shortest paths from all vertices to a destination vertex.
 - Like single-source shortest paths, simply invert all edges.
- **Single-pair shortest-paths problem (SPSP)**: find a shortest path between two vertices.
 - Actually not much easier than single-source shortest paths!
- **All-pairs shortest paths problem (APSP)**: find shortest paths between all pairs of vertices.
 - Trivial: solve single-source shortest paths for all vertices.
More clever solutions are more efficient.

► Dijkstra's algorithm for the SSSP


- **Idea from BFS:** Maintain **distance estimates .d** that are no smaller than shortest-path distances.
- Grow a set S of vertices whose **final shortest-path distances** from source s have been found.
每次维护每个最短路径的估计值
- **Idea from Prim:** In each step, **add the closest vertex** from $V \setminus S$ (smallest distance estimate .d → greedy choice).
- Refine distance estimates after each expansion of S .

DIJKSTRA(G, w, s)


```
1: Initialise  $d$  and  $\pi$  in the usual way.
2:  $S = \emptyset$ 
3:  $Q = V$ 
4: while  $Q \neq \emptyset$  do
5:    $u = \text{EXTRACT-MIN}(Q)$ 
6:    $S = S \cup \{u\}$ 
7:   for each  $v \in \text{Adj}[u]$  do
8:     if  $v.d > u.d + w(u, v)$  then
9:        $v.d = u.d + w(u, v)$ 
10:       $v.\pi = u$ 
11:      DECREASE-KEY( $Q, v, v.d$ )
```

每次选择离 S 最近的顶点，更新和他相邻的点


► Dijkstra's algorithm: Example




(a)




(b)




(c)



(d)




(e)



(f)

► Correctness of Dijkstra's algorithm

- We show that at the time a vertex u is added to S , $u.d$ is the shortest-path distance.
- This holds for s , so assume for a **contradiction** that $u \neq s$ is the **first vertex added to S** for which $u.d$ is larger than the shortest-path distance ($u.d \neq \delta(s, u)$).
- Consider a **shortest path p from s to u** and let y be the first vertex **outside of S** on this path. Let $x \in S$ be its predecessor.
- By choice of u , $x.d$ is the shortest-path distance to x , and when x was added, $y.d$ was set to $x.d + w(x, y) = \delta(s, y)$, the shortest-path distance to y (because otherwise **p** would not be the shortest to u).
- Since the path p_2 from y to u has non-negative distance,
$$y.d = \delta(s, y) \leq \delta(s, u) \leq u.d.$$
- Since u is added to S before y , $u.d \leq y.d$. Together $u.d = y.d$ and since y has the correct shortest-path distance, so has u , contradiction.



► Runtime of Dijkstra w/ Min-Heaps

- Runtime exclusive of red lines :
 $O(|V| + |E|)$
- Building a Min-Heap: $O(|V|)$.
- Runtime for all calls to Extract-Min is $O(|V|\log(|V|))$.
- Runtime for at most $|E|$ Decrease-Keys is $O(|E|\log(|V|))$.
- **Total:** $O((|V| + |E|) \log(|V|))$ or $O(|E| \log(|V|))$ if all vertices are reachable from the source.
- NB: for single-pair shortest paths we may stop when destination found.

DIJKSTRA(G, w, s)

1: Initialise d and π in the usual way.

2: $S = \emptyset$

3: $Q = V$

4: **BUILD-MIN-HEAP**(Q)

5: **while** $Q \neq \emptyset$ **do**

6: $u = \text{EXTRACT-MIN}(Q)$

7: $S = S \cup \{u\}$

8: **for each** $v \in \text{Adj}[u]$ **do**

9: **if** $v.d > u.d + w(u, v)$ **then**

10: **DECREASE-KEY**($Q, v.d, u.d + w(u, v)$)

11: $v.\pi = u$

► Summary

- Minimum spanning trees can be solved with two greedy algorithms:
 - Kruskal's algorithm adds the lightest edge connecting two trees
 - Prim's algorithm grows one tree by adding the lightest edge
- Dijkstra's algorithm solves single-source shortest paths by expanding on the set of vertices closest to the source.
 - Combines greedy and dynamic programming approaches
- Efficient data structures (union-find and priority queues) are vital for implementing the above algorithms efficiently.
- All algorithms can be implemented in time $O(|E|\log(|V|))$.
 - Advanced data structures (Fibonacci heaps) can improve this further.